



North Point Digital

The AI Adoption Playbook

A practical guide for small businesses: how to pick a use case worth proving, what AI really costs, and how to get working software instead of a strategy deck.

Read this first

Most of what's written about AI for business is either hype or homework. The hype promises everything and shows nothing. The homework reads like it was written for a bank's IT department. Neither is much use if you run a twenty-person company and want to know what to do on Monday morning.

Some context before anything else. UK government research published in February 2026 found that only **16% of UK businesses use any AI at all**. Whatever the headlines have made you feel, you are not behind. The gap between businesses that put AI to work and those that don't is opening now, though, and moving early is still cheap.

The other thing worth knowing is that when AI projects fail, it's almost never the technology. They fail earlier than that: wrong use case, unprepared data, nobody agreed what success meant, or the build quietly turned into a science project. Each section here goes after one of those failure points.

How to use this: read it once (twenty minutes), then work through the steps with your own business in mind. By the end you should have a candidate use case, a rough cost expectation, and a list of questions that will keep any provider honest — including us.

What's inside

- **Step 1** — Pick the job, not the technology
- **Step 2** — The data readiness check
- **The data layer** — Consolidate or connect · search · RAG, demystified
- **Step 3** — What it actually costs
- **Step 4** — Build it, buy it, or use no-code?
- **The landscape** — Everyday AI tools · agents & harnesses · workflow or agent · orchestration
- **Choosing the system** — The decision framework, and why the shape matters
- **Step 5** — Prove it in six weeks
- **Step 6** — Questions that keep providers honest

Pick the job, not the technology

Projects that start with "we should do something with AI" stall. Projects that start with "this task eats eleven hours a week" tend to work. Start from the job and you'll usually be fine. Start from the technology and you're shopping for a problem to justify it.

The best candidate jobs share three traits. They're **repetitive** (done weekly or daily). They're **rule-ish**, meaning a competent new hire could pick them up from examples. And the **information already exists** somewhere in your business, in an inbox or a folder or a system.

What this looks like in a real small business

Now	With AI
The inbox fills with the same twenty questions every week	Replies drafted from your past answers; a human checks and sends
Someone retypes invoices and delivery notes into the system	Amounts, dates and terms land in your accounts package on their own
The answer is in a manual or an old proposal — somewhere	Ask in plain English, get the answer with its source
Quotes start from a blank page	A draft estimate based on what similar jobs actually cost you
Call notes live in notebooks; the CRM is always behind	Meetings summarised and filed against the right customer
Every support ticket waits in one queue	Tickets sorted and routed, a suggested reply already drafted

The ten-minute exercise: ask your team what they retype, re-answer or look up again and again every week. Estimate the hours. Multiply by fifty. That number, in pounds, is the budget conversation. And the task at the top of the list? That's your candidate use case.

The data readiness check

AI is only as useful as the information you can put in front of it. Data readiness is the single biggest driver of cost and timeline in small-business AI projects — published industry analyses consistently find data preparation adds **40–60% to budgets** when nobody scoped it up front.

The good news: checking your own readiness takes an afternoon, not a consultancy engagement. For your candidate use case from Step 1, answer these honestly:

The five questions

- **1. Where does the information live?** One system is ideal. Two or three is normal. "Scattered across five tools and a shared drive" is workable but costs more.
- **2. Can you export it?** If you can produce a spreadsheet or a folder of documents today, you're ready to prove something. If exporting needs your IT supplier's help, get that conversation started now.
- **3. Is there enough history?** A few hundred past examples (emails answered, invoices processed, quotes written) is plenty for a proof of concept. Ten examples is not.
- **4. Is anything sensitive in there?** Customer personal data means GDPR applies — not a blocker, but it shapes where and how the AI runs. Flag it early.
- **5. Who knows the data best?** Name the person. A pilot without them assigned a few hours a week will drift.

Scoring it: mostly comfortable answers put you at the affordable end of any quote you'll receive. Two or more awkward answers doesn't mean stop — it means the first week of any engagement should be data scoping, and a provider who skips that step is guessing with your money.

And one myth worth retiring. You don't need a data warehouse to start, whatever you've been told. You don't need a particular cloud either, or anyone technical on staff. Exports and email will prove a use case; the infrastructure conversation belongs after the proof, not before it.

Consolidate or connect?

The most expensive sentence in small-business AI is: *"First we need to get all our data in one place."* It sounds responsible. What actually happens is a six-month tidying project that delays every AI benefit and then quietly stalls, because nobody is waiting for tidy data. They're waiting for results.

The modern answer is mostly the opposite: **leave data where it lives and connect to it**. AI systems read from the tools you already use, through connectors and exports. Your accounts stay in the accounts package. The AI visits.

The three honest options

Approach	When it's right
Connect in place	The default. Connectors or scheduled exports from each system. Days of work, not months — and if a use case dies, nothing was wasted
Light consolidation	For documents: one well-named folder structure for the files that matter (policies, proposals, manuals). An afternoon's discipline that pays forever
Full warehouse	When reporting across systems is itself the goal, or volumes are genuinely large. Rarely the right first move for a small business — and never a prerequisite for proving an AI use case

Access: who — and what — can see your data

- **Least privilege.** The AI should see only what its task needs. An invoice reader needs the invoices folder, not the HR drive.
- **Service accounts, not personal logins.** Access granted to the system under its own identity — revocable, auditable, and not tied to whoever set it up.
- **An audit trail.** You should be able to answer "what did it read, and when?" If a provider can't show you that, see Step 6.
- **Personal data means GDPR.** Not a blocker — but where the processing happens, and under what agreement, must be a written answer.

The rule: never let a data project block an AI project. Prove the use case on an export, and consolidate only what the proven use case demands. You'll find the proof tells you which data was actually worth tidying.

Searching everything you know

Every business more than a few years old has the same problem. The answer exists, somewhere in an old email or a proposal or someone's notebook, but finding it costs more than redoing the work. Fixing this is probably the most boring high-value AI project there is, and it's now well within a small business's reach.

Why AI search is different


Traditional search matches words: search "staff handbook", miss the file called "employee manual". AI search matches **meaning** — the technique is called semantic (or vector) search, and it's why a question phrased your way finds a document phrased someone else's way.

The second piece is the one you'll hear at every vendor meeting: **RAG**, short for retrieval-augmented generation. In one line, the AI looks things up first, then answers from what it found and cites its sources. Vendors love the acronym, so the next page takes it apart properly.

What it takes to work

- **Connectors** to where knowledge lives — shared drives, email, the CRM, the support desk — with the index kept up to date automatically.
- **Permission-aware results.** The answer must respect who's asking: the sales team's search should never surface HR documents. Insist on this; it's where cheap tools cut corners.
- **One good corpus first.** Start with the best-kept set of documents — policies, manuals, past proposals — not the whole chaotic drive on day one.


The honest limitation: search quality mirrors document hygiene. If three versions of a price list are indexed, the AI will confidently cite the wrong one. The fix is the light consolidation from the previous page — retire duplicates and outdated files from the corpus before you index it, and add to it gradually.



RAG, demystified

Here's the problem RAG solves. A model knows nothing about *your* business, and when you ask it anyway, it improvises. Fluently, confidently, and wrongly. Retraining a model on your data costs real money and goes stale the day a document changes. RAG sidesteps both: it's an open-book exam instead of a recital.

What actually happens, in four steps



What this buys you: answers **stay current**, because changing the document changes the next answer with no retraining. And answers are **checkable**, because the citation is your defence against confident nonsense.

What's changed lately — the advancements, demystified

You'll hear	What it actually means
Hybrid search	Meaning-based and exact-keyword search combined — meaning alone fumbles part numbers and names. Table stakes now
Reranking	A second pass re-orders retrieved passages so the best reach the model. Cheap, big accuracy win
Graph-based RAG	Maps how entities relate across documents, so it can answer "what themes run across all our complaints?" — the whole, not a passage
Agentic retrieval	Searches like a person: rephrases, follows leads, checks, searches again. Better on hard questions; slower and dearer
Long context instead	Frontier models read huge amounts in one go. If your corpus is one manual, pasting it in beats building a pipeline — an honest vendor will say when you don't need RAG at all

Vendor questions that sort the field: Hybrid search? Reranker? Citations on every answer? Permission-aware results? How fast does the index pick up a changed document? Two minutes, and you'll know more than most buyers.

What it actually costs

Almost nobody in this industry publishes prices, which makes budgeting feel like guesswork. It isn't. These are the real UK price bands as of 2026, from the published price guides and from our own quoting:

Engagement	Typical UK price	What you get
Feasibility check	£5k–£15k	4–6 weeks, one narrow question answered, usually throwaway code
Proof of concept	£15k–£85k	6–12 weeks, working software on your data, production costings
Production build	£75k+	The full system, integrated and supported — only after a proof

Behind the spread sit day rates. UK freelance AI practitioners charge roughly £400–£800 a day, specialist consultancies £1,200–£2,500, and the big firms more again. A normal proof of concept is twenty to thirty senior days. Multiply through and the bands stop looking arbitrary.

The costs nobody quotes

- **Running costs.** A pilot costs tens to low hundreds of pounds a month to run. Production can differ sharply. Any proposal without a running-cost estimate is incomplete.
- **Data preparation.** The 40–60% surprise from Step 2 — make sure it's inside the quote, not waiting behind it.
- **Your own time.** Plan for a few hours a week from the person who knows the work. Pilots that get no internal attention quietly die.

Insist on a fixed fee. A proof of concept is a defined question with an end date; an open-ended day rate hands all the risk to you. Roughly half of UK SME engagements are now fixed-price — a provider who won't fix a price for their own recommended scope is telling you something.

Worth checking before you assume the headline price: Innovate UK has run AI proof-of-concept grant rounds funding up to 70% of eligible costs for micro and small businesses on larger projects. Eligibility shifts round to round, but ten minutes of checking can change the maths.

Build it, buy it, or use no-code?

Not every use case needs a consultancy, and you should be suspicious of any guide that pretends otherwise. So, the decision tree:

Off-the-shelf is enough when...

The job is generic — drafting documents, summarising meetings, general research. ChatGPT, Claude or Copilot at £15–25 per person per month covers it. Train your team on it (Step 1's exercise works here too) and bank the easy win. No project required.

No-code tools are enough when...

The job is a simple, low-stakes workflow. A basic chatbot answering from a single document, say, or a form that triggers a templated email. The current crop of no-code AI builders genuinely handles this, cheaply. Their limit arrives the moment a prototype becomes *load-bearing*, when real customers and real edge cases start hitting it. If a toy fails quietly you've learned something. If your quoting process fails quietly, you have a different sort of problem.

Bespoke (and help) earns its keep when...

- The AI must read **your** history — your answers, your pricing, your documents — not general knowledge
- It must write into systems you rely on (accounts, CRM, support desk)
- Errors carry a real cost: wrong prices quoted, wrong customers emailed, compliance exposure
- Customer personal data is involved and "we pasted it into a chatbot" won't survive a GDPR question

The pattern that works: start everyone on off-the-shelf tools this month. It costs little and builds intuition. In parallel, take the one use case with real money attached through a structured proof. The first habit gets your people comfortable; the second is what actually changes the business.

The assistants your team should already be using

Before any project, there's the £20-a-month tier: general-purpose AI assistants. As of mid-2026, four frontier models sit close together at the top — closer than the marketing suggests — and each has a personality worth knowing:

Tool	Worth knowing
Claude (Anthropic)	Consistently strong on writing and coding; the most natural-sounding text and a steady voice across long documents — useful for customer-facing copy
ChatGPT (OpenAI, GPT-5.5)	The safe generalist: neck-and-neck with Claude on coding, a favourite for creative work
Gemini (Google, 3.1 Pro)	Leads on reasoning and reads enormous amounts of text in one go — hand it the whole document pile. Also the cheapest frontier option per call
Microsoft Copilot	Less a model than a delivery mechanism: frontier AI inside Word, Excel, Outlook and Teams. If your business lives in Microsoft 365, it's the lowest-friction starting point

Notice what's missing from that table: a winner. There is **no "best" model**, and the rankings reshuffle every few months. That's good news for you. Prices keep falling, and there's no reason to marry a vendor.

The two rules that outlast every release cycle: test on your own data, because a benchmark says nothing about how a model handles *your* invoices and *your* customers; and keep whatever you build swappable, so the underlying model can change without starting again.

Agents and harnesses, in plain English

A chatbot answers. An agent acts. That's the whole distinction. An agent doesn't just tell you what to write back to a customer — it reads the message, looks up the order, drafts the reply and files the record, then asks a human to approve. The examples in Step 1 are all agent work.

Agents are also genuinely early. Gartner expects around 40% of enterprise applications to include task-specific agents by the end of 2026, up from fewer than 5% in 2025. Deloitte, meanwhile, found only about a fifth of companies have any mature way of governing them. In other words: real momentum, with the guardrails lagging behind it. Both halves of that are worth taking seriously.

The harness: the part that does the work


The model gets the headlines, but a model on its own is a very capable text box. What turns it into something that can do a job is the layer wrapped around it — the **harness**: the tools the agent can use, the memory it keeps, and the guardrails, approvals and audit trail that make it trustworthy. If the model is the engine, the harness is the rest of the car.

- **Hermes** (Nous Research) — an open-source harness that wraps around any model and gives it tools, layered memory and long-running, multi-agent structure. The thesis: a small team with open tooling can stand up an agent rivalling commercial offerings, without handing data to a single vendor.
- **OpenClaw** — a self-hosted, MIT-licensed gateway connecting the chat apps your team already uses (WhatsApp, Slack, Teams, iMessage) to an agent you run yourself. Its quiet superpower is adoption: nobody has to learn a new app.
- **Coding agents** like Claude Code sit in the same family — agents with a harness purpose-built for software work, often driven from the tools above.

The rule that matters more than the tooling: keep a person at the decision point. Let the agent do the legwork; let a human approve anything that touches money, customers or reputation. Every stable small-business agent we've seen follows it.

Workflow or agent?

These two words get used interchangeably in sales conversations, and they shouldn't be. They're opposite philosophies, and they fail in different ways. A **workflow** is a route you've drawn. An **agent** is a destination you've set.



	Workflow	Agent
Behaviour	Predictable — same input, same route	Adaptive — chooses its route per case
Best at	Repetitive, well-understood processes	Variation, judgement, messy inputs
Running cost	Low and stable	Higher, and varies with the task
When it breaks	Loudly, at the edge cases	Quietly, in ways you must watch for
Audit	Trivial — the route is the documentation	Needs deliberate logging and review

The middle path most small businesses actually need

The highest-value pattern in small-business AI is neither extreme. It's a **workflow with AI steps**: a deterministic skeleton, fixed route, auditable and cheap, with AI handling the one step rules can't manage, like reading the document or drafting the reply. You get judgement where it's needed and predictability everywhere else.

To map the examples from Step 1: invoice retyping is a workflow with one AI extraction step. So is quote drafting. "Ask your documents" is search (see The data layer), not an agent at all. Inbox triage with judgement calls across systems, though — *that's* agent territory, with a human approving the sends.

Vendor test: ask "could this be a workflow with one AI step?" If the seller of an agent platform can't answer crisply, they're selling you their architecture, not your solution.

One agent or many?

Once one agent works, the temptation is to wire up a whole team of them. Sometimes that's the breakthrough; often it's complexity you didn't need. A **single agent** — one model in a loop with good tools — is simpler, cheaper, far easier to debug, and handles more than people expect. Start there.

The tools you'll hear named

Tool	What it is, honestly
n8n	Workflow automation with AI steps — visual flows connecting your apps, with model calls in the middle. The most small-business-friendly entry point: many "agent" needs are really an n8n workflow with one clever step
CrewAI	Role-based agent "crews" modelled on a human team. Fast to prototype; its convenient abstractions can get in the way at production scale, and observability is weaker
LangGraph	Work modelled as a graph of steps with state, checkpoints and human-in-the-loop approvals. The sensible default when a workflow needs branching or auditability
Vendor SDKs	The OpenAI Agents SDK and Anthropic's Claude Agent SDK — often the quickest path for a single agent with a tool or two, no framework overhead

When many beats one — and when it doesn't

Multi-agent earns its keep when the work genuinely splits into roles with different tools and the volume justifies the overhead. It is **not** worth it while you're still experimenting — or, most importantly, while you can't yet reliably observe and debug *one* agent. Adding agents to a system you can't see into doesn't add capability; it adds places for things to go wrong quietly.

For most small businesses the sequence is: assistants first, then one agent on one workflow with a human approval step, and orchestration only if the numbers demand it. Skip steps and you'll end up with a great demo and not much else.

The decision framework

Six questions decide the right shape for any automation. Answer them for your candidate use case before any vendor conversation, and you'll walk in knowing what you need:

Question	What the answer points to
1. Could you write the steps on one page?	Yes → workflow. No, "it depends" every time → agent territory
2. Is any step reading or writing natural language?	That step is where AI goes — the rest can stay plain automation
3. What does a wrong action cost?	If it's real money or reputation: human approval before anything irreversible, whatever the shape
4. How many times a week does this happen?	A handful → an assistant and a person may be enough. Dozens+ → automation pays
5. Does each case need judgement, or just consistency?	Consistency → workflow. Judgement → AI step or agent
6. Could you observe and debug one agent today?	If not, you're not ready for several — build the simpler shape first

The ladder, and why you climb it slowly



Climb a rung only when the one below demonstrably fails the job. Why? Because the wrong shape costs you in both directions. An agent doing a workflow's job means paying, in money and supervision, for flexibility nobody asked for. A workflow doing an agent's job means brittle rules sprouting exceptions until no one understands the rule garden any more. Both failures start the same way: someone chose the system before they understood the job.

If in doubt: build the workflow, put AI in one step, keep a human at the end. It's the cheapest shape to be wrong about — you'll learn more from two weeks of it running than from any architecture debate.

Prove it in six weeks

However you source the work, the proof phase should follow the same shape. Six weeks is enough to prove or disprove almost any small-business use case. If an engagement runs longer than that without showing working software, it's drifting.

The shape



The rules that keep it honest

- **One use case.** Not a platform. The narrower the question, the cheaper the answer.
- **Your data from week one.** A demo on sample data proves the vendor can demo. Only your data proves anything about your business.
- **Agree the kill criteria up front.** Decide before building what accuracy, speed or adoption number means "don't proceed". A pilot that can't fail is a brochure.
- **Weekly demos, no exceptions.** If you can't see progress, you're funding a surprise.
- **"Don't proceed" is a result.** A documented no after six weeks beats eighteen months of maybe, and costs a tenth as much.

What you should hold at the end — whoever does the work: software you can put in front of users, honest production running costs, a written route to live, and ownership of the code and prompts. A slide deck with a roadmap is not a proof of concept, whatever the invoice says.

Questions that keep providers honest

Take these into any conversation, including one with us. Good providers answer them quickly and in writing. The other kind change the subject.

- **"What's the fixed price for the scope you're recommending?"** Watch for day rates dressed as estimates.
- **"What will it cost to run in production, monthly?"** If the thinking hasn't been done, the proposal is incomplete.
- **"What happens if the pilot shows it doesn't work?"** The honest answer involves the phrase "we tell you to stop".
- **"Who owns the code, prompts and any trained models?"** The answer should be: you. Otherwise you're renting, and the price should say so.
- **"Where does our data go, and who can see it?"** You want named infrastructure and a straight GDPR answer, not reassurance.
- **"Who exactly does the work?"** Meet the engineers, not just the salesperson. Ask whether the people on the call are the people on the project.
- **"Can I speak to a past client?"** One reference call tells you more than any deck.
- **"What do you need from us, and how many hours a week?"** A provider who says "nothing" hasn't done this before.

Five ways projects stall (so you can spot them early)

- The use case was chosen in a boardroom, not from the work itself
- Data scoping was skipped and surfaced later as cost and delay
- Success was never defined, so the pilot could neither succeed nor fail
- The prototype impressed everyone and integrated with nothing
- No one inside the business owned it, so it ended with the invoice

Who we are, and the offer

North Point Digital is a UK-based team of AWS-certified architects and engineers. We spent the last decade inside enterprise cloud estates, building them, migrating them, and finding out where the money goes. These days we point that engineering at AI for small businesses.

Everything this playbook recommends is how we actually work:

- **The AI Launchpad** — a six-week proof of concept exactly as described in Step 5: your data, weekly demos, working software, honest costings, and a documented recommendation either way. Typically **£15,000–£25,000, fixed**, agreed in writing before we start. You don't need to be technical, or on any particular cloud.
- **The Security & Cost Optimisation Blueprint** — for AWS customers: a two-week review targeting 25%+ savings, with a written guarantee — if we don't find at least 25% in potential savings, you don't pay. Across 20+ reviews, the average found is 30%.

No long-term contracts on anything. The first call is thirty minutes, free, with an engineer — and if there's no fit, we'll say so on the call.

Book a free call: northpointdigital.com — or email contact@northpointdigital.com and tell us about the task at the top of your Step 1 list.

Sources for figures used in this guide: UK government AI adoption research (Feb 2026); published UK AI consultancy price guides (2026); industry analyses of AI project cost overruns; Innovate UK grant programme documentation. Detailed citations at northpointdigital.com/blog.